

Purabh singh - 1nt22is122 Palak kota - 1nt22is110 Prathuysa NB - 1nt22is119 Monika K - 1nt22is097

```
In [1]: import pandas as pd
from xgboost import XGBClassifier
from sklearn.model_selection import train_test_split
from sklearn.metrics import classification_report, confusion_matrix
from sklearn.utils.class_weight import compute_sample_weight
import joblib

In [2]: # Load the final features with labels
df = pd.read_csv("final_labeled_features.csv")

# Convert multi-class labels to binary:
# 1 = has symptoms, 0 = no symptoms
df['label'] = df['label'].replace({
    'no symptoms': 0,
    'stress': 1,
    'anxiety': 1,
    'depression': 1
})

# Drop non-numeric columns (text and patient ID)
X = df.drop(columns=['label', 'utterance', 'patient_id'], errors='ignore')
y = df['label']

In [3]: X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, stratify=y, random_state=42
)

In [4]: # Compute weights for each sample to balance classes
sample_weights = compute_sample_weight(class_weight='balanced', y=y_train)

In [26]: xgb_model = XGBClassifier(
    n_estimators=200,
    learning_rate=0.1,
    max_depth=6,
    objective='binary:logistic',
    eval_metric='logloss',
    use_label_encoder=False,
    random_state=42
)

xgb_model.fit(X_train, y_train, sample_weight=sample_weights)
print(xgb_model.get_params())

/Users/Saravanan/opt/anaconda3/lib/python3.9/site-packages/xgboost/core.py:158: UserWarning: [21:17:27] WARNING: /Users/runner/work/xgboost/xgboost/src/learner.cc:740:
Parameters: { "use_label_encoder" } are not used.

warnings.warn(msgs, UserWarning)
{'objective': 'binary:logistic', 'base_score': None, 'booster': None, 'callbacks': None, 'colsample_bylevel': None, 'colsample_bynode': None, 'colsample_bytree': None, 'device': None, 'early_stopping_round': None, 'enable_categorical': False, 'eval_metric': 'logloss', 'feature_types': None, 'gamma': None, 'grow_policy': None, 'importance_type': None, 'interaction_constraints': None, 'learning_rate': 0.1, 'max_ax_bin': None, 'max_cat_threshold': None, 'max_cat_to_onehot': None, 'max_delta_step': None, 'max_depth': 6, 'max_leaves': None, 'min_child_weight': None, 'missing': nan, 'monotone_constraints': None, 'multi_strategy': None, 'n_estimators': 200, 'n_jobs': None, 'num_parallel_tree': None, 'random_state': 42, 'reg_alpha': None, 'reg_lambda': None, 'sampling_method': None, 'scale_pos_weight': None, 'subsample': None, 'tree_method': None, 'validate_parameters': None, 'verbosity': None, 'use_label_encoder': False}

above are the available parameters of the respective model. And Only a few key parameters are usually set manually. that is { 'n_estimators': 200, 'learning_rate': 0.1, 'max_depth': 6, 'objective': 'binary:logistic', 'use_label_encoder': False, 'random_state': 42, 'eval_metric': 'logloss', ... }

In [6]: y_pred = xgb_model.predict(X_test)

print("\n Classification Report:")
print(classification_report(y_test, y_pred, digits=3))

print("\n Confusion Matrix:")
print(confusion_matrix(y_test, y_pred))
```


		precision	recall	f1-score	support
	0	0.781	0.435	0.559	115
	1	0.926	0.983	0.954	828
	accuracy			0.916	943
	macro avg	0.854	0.709	0.756	943
	weighted avg	0.908	0.916	0.906	943
	Confusion Matrix:				
	[[50 65]				
	[14 814]]				

```
In [7]: joblib.dump(xgb_model, "xgboost_depression_model.pkl")
print(f"✔ Model saved as xgboost_depression_model.pkl")

✔ Model saved as xgboost_depression_model.pkl
```

```
In [8]: import matplotlib.pyplot as plt
import xgboost as xgb

xgb.plot_importance(xgb_model, max_num_features=15, importance_type='gain')
plt.title("Top 15 Feature Importances")
plt.show()
```



```
In [31]: from sklearn.metrics import accuracy_score

accuracy = accuracy_score(y_test, y_pred)
print(f"✔ Overall Accuracy: {accuracy:.4f}")

✔ Overall Accuracy: 0.9162
```

```
In [30]: from sklearn.metrics import classification_report, confusion_matrix, accuracy_score

# Predictions
y_pred = xgb_model.predict(X_test)

# Evaluation
print("\n Classification Report:")
print(classification_report(y_test, y_pred, digits=3))

print("\n Confusion Matrix:")
print(confusion_matrix(y_test, y_pred))

# Overall Accuracy
accuracy = accuracy_score(y_test, y_pred)
print(f"✔ Overall Accuracy: {accuracy:.4f}")
```

20/05/2025, 21:40labeled_model_final

Classification Report:

	precision	recall	f1-score	support
0	0.781	0.435	0.559	115
1	0.926	0.983	0.954	828
accuracy			0.916	943
macro avg	0.854	0.709	0.756	943
weighted avg	0.908	0.916	0.906	943

Confusion Matrix:

[[50 65]

[14 814]

Overall Accuracy: 0.9162

In []:

In [21]:

```
from sklearn.metrics import roc_curve, roc_auc_score
import matplotlib.pyplot as plt
```

In [22]:

```
# Get probabilities for the positive class (class 1 = symptoms)
y_proba = xgb_model.predict_proba(X_test)[:, 1]
```

In [27]:

```
# Compute FPR, TPR, and thresholds
fpr, tpr, thresholds = roc_curve(y_test, y_proba)

# Compute AUC
auc_score = roc_auc_score(y_test, y_proba)
print(f" ROC-AUC Score: {auc_score:.4f}")

ROC-AUC Score: 0.9220
```

In [24]:

```
plt.figure(figsize=(8, 6))
plt.plot(fpr, tpr, label=f"AUC = {auc_score:.4f}", color='darkorange')
plt.plot([0, 1], [0, 1], linestyle='--', color='gray') # Diagonal line
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve for Symptom Detection')
plt.legend(loc='lower right')
plt.grid(True)
plt.show()
```

The figure is a Receiver Operating Characteristic (ROC) curve titled "ROC Curve for Symptom Detection". The x-axis is labeled "False Positive Rate" and ranges from 0.0 to 1.0. The y-axis is labeled "True Positive Rate" and also ranges from 0.0 to 1.0. A dashed gray diagonal line represents the performance of a random classifier. An orange solid line represents the performance of the machine learning model. The orange curve starts at (0,0), rises steeply to a True Positive Rate of approximately 0.7 at a False Positive Rate of 0.05, then continues to rise more gradually, reaching a True Positive Rate of 1.0 at a False Positive Rate of approximately 0.8. A legend in the bottom right corner indicates "AUC = 0.9220".

The plot will show:

TPR (True Positive Rate) vs. FPR (False Positive Rate)

AUC score tells how well your model distinguishes between the two classes:

localhost:8888/nbconvert/html/patient data/labeled_model_final.ipynb?download=false

3/5

20/05/2025, 21:40labeled_model_final

0.5 = random guessing

> 0.9 = excellent model

Title:

Detection of Stress, Anxiety, and Depression from Transcripts and Voice Data using Machine Learning

Objective:

To develop an ML model that predicts whether a patient exhibits symptoms of mental health issues (stress, anxiety, depression) based on both **textual transcripts** and **acoustic voice features**, using **unlabeled datasets** and zero-shot classification.

Dataset Description:

Each patient has three types of files:

1. **FORMANT.csv** — 5 formant frequency values per row (speech-related acoustic features).

2. **COVAREP.csv** — Rich set of low-level descriptors (voice features).

3. **TRANSCRIPT.csv** — Contains timestamped speech by 'Ellie' (AI interviewer) and the 'Participant'.

Example format:

335_FORMANT.csv

335_COVAREP.csv

335_TRANSCRIPT.csv

PHASE 1: Preprocessing and Feature Engineering

1. Transcript Processing

Extracted utterances spoken by the `Participant`.

Synced utterances to their `start_time` and `stop_time`.

2. Audio Feature Extraction

Extracted **Formant** and **COVAREP** values aligned to each utterance's timestamp.

Calculated **mean**, **std**, and **range** for each segment.

3. Text Embedding

Used **all-MiniLM-L6-v2** model from `sentence-transformers` to convert utterances into 384-dimensional vector embeddings.

4. Feature Merging

Combined all features (formant + covarep + text embedding) into a single flat vector.

Created a `DataFrame` where each row = one utterance's full feature set.

5. Zero-Shot Labeling

Used Hugging Face's `facebook/bart-large-mnli` to classify each utterance as:

`stress`, `depression`, `anxiety`, or `no symptoms`.

Mapped all mental-health-related labels into a single class: `1 = symptoms`, `0 = no symptoms`.

6. Final dataset: `final_labeled_features.csv`

~4,700 utterances × ~860 features + binary label.

PHASE 2: Model Training

localhost:8888/nbconvert/html/patient data/labeled_model_final.ipynb?download=false

4/5

1. **Model Chosen:**
- **XGBoost (XGBClassifier)** — selected for its performance on tabular, high-dimensional, and imbalanced data.
2. **Data Split:**
- 80/20 train-test split (stratified by label).
3. **Class Imbalance Handling:**
- Used `compute_sample_weight` to address class imbalance during training.
4. **Training Parameters:**
- ```
XGBClassifier(
 n_estimators=200,
 learning_rate=0.1,
 max_depth=6,
 objective='binary:logistic',
 eval_metric='logloss',
 use_label_encoder=False,
 random_state=42
)
```
5. Model saved: `xgboost_depression_model.pkl`

PHASE 3: Evaluation

1. **Metrics Used:**
- Accuracy
  - Precision, Recall, F1-Score
  - Confusion Matrix
  - ROC Curve & AUC Score
2. **Sample Output:**
- ```
Classification Report:  
              precision    recall  f1-score  
no symptoms      0.74      0.42      0.53  
symptoms         0.92      0.98      0.95  
Accuracy: 0.9118  
AUC Score: 0.94
```
3. **ROC Curve:** Visualized using `matplotlib`, clearly showed strong separation between the two classes.

🧰 Tools & Libraries Used

- Python, Pandas, NumPy
- XGBoost
- Scikit-learn
- SentenceTransformers
- Hugging Face Transformers
- Matplotlib, Seaborn (for visualizations)
- Joblib (for model saving)

In []: